Course Project 4
Backreferences and NP-Completeness

CSE 30151 Spring 2018

Version of January 25, 2018

In this project, we return to `grep` and allow backreferences inside regular expressions themselves. For example, in the regular expression

$$((a|b)*)\backslash 1$$

the backreference (`\1`) must match a substring equal to the contents of group 1. Thus this expression recognizes the language $\{ww \mid w \in \{a,b\}^*\}$. This new feature takes `grep` beyond regular languages and even beyond context-free languages. In fact, you'll show that this makes `grep` NP-complete by implementing a SAT solver with it.

You will need a correct solution for CP3 to complete this project. If your CP3 doesn't work correctly (or you just weren't happy with it), you may use the official solution or another team's solution, as long as you properly cite your source.

Getting started

To make sure your repository is up to date, please have one team member run the commands

```
git pull https://github.com/ND-CSE-30151-SP18/theory-project-skeleton
git push
```

and then other team members should run `git pull`. The project repository should then include the following files:

```
bin/
  parse_re
  bgrep
  sat
tests/
  test-cp4.sh
cp4/
```

Please place the programs that you write into the `cp4/` subdirectory.

1 Backreferences: Syntax

Extend your regular expression parser to allow backreferences. The syntax of backreferences is the same as in the replacement strings of CP3. In terms of the context-free grammar of CP2, you should add new rules

$$\text{Primary} \rightarrow \backslash \text{ Number}$$
$$\text{Primary} \rightarrow \backslash \text{ g < Number >}$$

where Number rewrites to any positive integer (without leading zeros). Here are some example regular expressions with their abstract syntax trees:

expression	syntax tree
\1	(backref 1)
\01	error
\10	(backref 10)
\10*	(star (backref 10))
\g<10>	(backref 10)
\g<1>0	(concat (backref 1) (symbol "0"))

To test your extension, update `parse_re` from CP2 to handle backreferences. For example:

```
$ parse_re '(a)\1\g<1>'
(concat (group (symbol "a")) (backref 1) (backref 1))
```

Test by running `test-cp4.sh`.

2 Backreferences: Semantics

To understand backreferences in more detail, consider the expression

$$(a|b)*\backslash 1$$

Note that, unlike the example at the beginning of this document, the star is now outside the group. If group 1 matches multiple substrings, then, just as in CP3, it's only the last substring whose contents are used. Thus `abb` matches but `aba` does not.

On the other hand, if group 1 does not match any substrings, then nothing (not even the empty string) can match \1. So, if the input string is ε, group 1 does not match any substring. So \1 cannot match anything, so the input string does not match.

Here's another example:

$$(aaa*)\backslash 1\backslash 1*$$

This recognizes the language $\{\mathtt{a}^n \mid n \text{ is not prime}\}$. Do you see why?

There's no speed requirement in this project, so you have considerable latitude in how you implement backreferences. You're not even required to continue using NFAs, but if you do, you can convert each backreference $\backslash k$ into the following NFA:

where "copy k" is a special transition like "open k" and "close k" from CP3.

In CP1, you wrote a function to test whether a NFA M accepts a string w using a search through the graph of configurations (q, i), where $q \in Q$ and $0 \le i \le |w|$. Suppose that the current configuration is (q, i) and the contents of group k is g_k. If there is a transition $q \xrightarrow{\text{copy } k} r$, we can check whether $w_i \cdots w_{i+|g_k|-1} = g_k$. If so, we can create a new configuration $(r, i + |g_k|)$. But to do this, we need to be able to know what g_k is, and our configurations don't contain this information. So, we can redefine configurations to include information about where groups have been opened and closed so far.

For example, consider again the regular expression $(\mathtt{aaa*})\backslash 1\backslash 1*$. This is equivalent to the NFA shown in Figure 1. On string $\mathtt{aaaaaa}$, the graph of configurations is shown in Figure 2. Each node is a configuration, which now includes not only a state and string position, but also information about the start and end of each group if it is known. Note that there are now two configurations for state q_6 and position 6, because there are two ways to get there that have two different contents for group 1.

Modify your NFA matching function to handle backreferences. You can use the method sketched above, or some other method. Depending on how much information you put into configurations, you may or may not need to watch out that regular expressions like $\mathtt{()*}$ don't cause an infinite loop. (The test script includes a test for this, so if the tests don't hang, then there's nothing to worry about.)

Finally, write a program called $\mathtt{bgrep}$ (backtracking $\mathtt{grep}$) that has the same usage as $\mathtt{agrep}$ but allows backreferences inside regular expressions, as described above. Run $\mathtt{test\text{-}cp4.sh}$ to test it.

3 SAT solver

Adding backreferences to regular expressions increases their power a lot; in fact, it makes matching NP-complete. Write a program that demonstrates this by reducing Boolean satisfiability to regular expression matching with backreferences.

$\mathtt{cnf\_to\_re}$ $\mathtt{formula}$

- $\mathtt{formula}$: formula in conjunctive normal form (see below)

- Output:

Figure 1: NFA equivalent to the regular expression `(aaa*)\1\1*`. Useless epsilon transitions are omitted for simplicity.
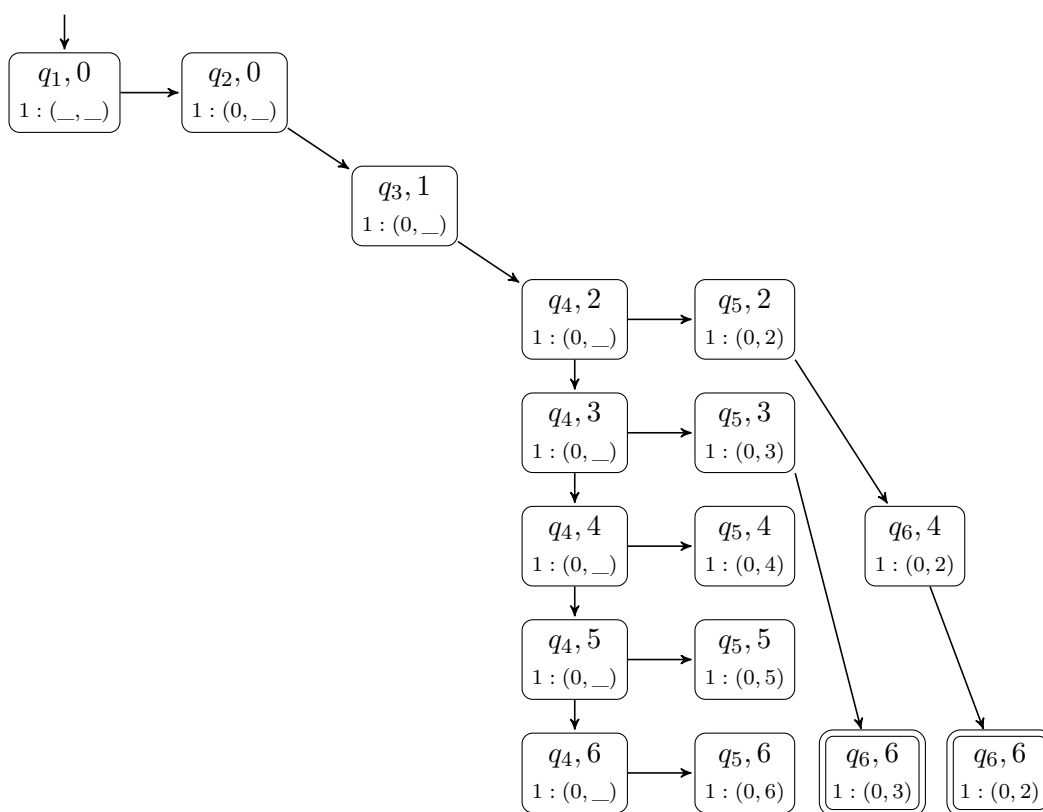

Figure 2: Graph of configuration for the NFA of Figure 1 and the input string `aaaaaa`.

4

- **regexp:** immediately (without intervening whitespace) followed by a regular expression α
- **string:** immediately (without intervening whitespace) followed by a string w

The string w should match the regular expression α if and only if the `formula` is satisfiable.

The `formula` should use &, |, and ! for logical and, or, not, respectively. Assume that the `formula` is in conjunctive normal form, and that each clause is surrounded by parentheses.

For example, if the formula is $(x_1 \vee x_1 \vee x_2) \wedge (\overline{x_1} \vee \overline{x_2} \vee \overline{x_2}) \wedge (\overline{x_1} \vee x_2 \vee x_2)$ is, then you would run the program like this:

```
$ cnf_to_re '(x1|x1|x2)&(!x1|!x2|!x2)&(!x1|x2|x2)'
regexp:...
string:...
```

Run `test-cp4.sh` to test. The test script also measures the running time of your program, which must be polynomial.

Submission instructions

Your code should build and run on studentnn.`cse.nd.edu`. The automatic tester will clone your repository, run `make -C cp4`, and then run `tests/test-cp4.sh`. You're advised to try all of the above steps and ensure that all tests pass.

To submit your work, please push your repository to Github and then create a new release with tag version `cp4` (note that the tag version is not the same thing as the release title). If you are making a partial submission, then use tag version `cp4-1` for part 1, `cp4-2` for part 2, and so on.

Rubric

Part 1 (`parse_re`)	3
Part 2	
matching backreferences	6
group matches multiple times	3
group matches no times	3
`bgrep`	3
Part 3	
reading CNF formulas	3
correctness	6
polynomial time	3
Total	30